



Parlib API Reference

Release 1.0

Kevin Klues, Barret Rhoden

February 07, 2013

CONTENTS

1	Welcome to Parlib’s API Reference!	1
2	API Reference	3
2.1	Vcores	3
2.2	Uthreads	5
2.3	MCS Locks	6
2.4	Spinlocks	8
2.5	Dynamic Thread Local Storage	8
2.6	Thread Local Storage	9
2.7	Memory Pools	9
2.8	Slab Memory Allocator	10
2.9	Atomic Memory Operations	11
3	About	13
3.1	People	13
3.2	Publications	13
	Index	15

WELCOME TO PARLIB'S API REFERENCE!

Parlib is a library meant to ease the development of software written for highly parallel systems. It was originally written for the Akaros operating system, and has since been ported to Linux. It was originally designed as an emulation layer on top of Linux that would allow developers to write applications and test them on a linux system before deploying them on Akaros. Since then, however, Parlib has proved itself useful in its own right as a standalone library for Linux. Most notably, as the backend for the Lithe implementation.

The official Home Page for Parlib is: <http://akaros.cs.berkeley.edu/parlib>.

At present, Parlib provides 6 primary services to developers:

1. A user-level abstraction for managing physical cores in a virtualized-namespaces (vcores)
2. A user-level thread abstraction (uthreads)
3. User-space MCS locks, MCS barriers, spinlocks and spinbarriers
4. Static and dynamic user-level thread local storage
5. Pool and slab allocator dynamic memory management abstractions
6. A standardized API for common atomic memory operations

API REFERENCE

2.1 Vcores

The vcore abstraction gives user-space an abstraction of a physical core (within a virtual namespace) on top of which it can schedule its user-level threads. When running code associated with a vcore, we speak of being in **vc core context**.

Vcore context is analogous to the interrupt context in a traditional OS. Vcore context consists of a stack, a TLS, and a set of registers, much like a basic thread. It is the context in which a user-level scheduler makes its decisions and handles signals. Once a process leaves vcore context, usually by starting a user-level thread, any subsequent entrances to vcore context will be at the top of the stack at `vc core_entry()`.

The vcore abstraction allows an application to manage its cores, schedule its threads, and get the best performance it possibly can from the hardware.

To access the vcore API, include the following header file:

```
#include <parlib/vcore.h>
```

2.1.1 Constants

MAX_VCORES

The maximum number of vc ores supported on this platform

LOG2_MAX_VCORES

The log-base-2 of the maximum number of vc ores supported on this platform

2.1.2 Types

```
struct vc core  
vc core_t
```

An opaque type used to maintain state associated with a vcore

2.1.3 External Symbols

```
extern void vc core_entry()
```

User defined entry point for each vcore. If `vc core_saved_ucontext` is set, this function should just restore it, otherwise, it is user defined.

2.1.4 Global Variables

`vcore_t *__vcores`

An array of all of the vcores available to this application

`void **__vcore_tls_descs`

An array of pointers to the TLS descriptor for each vcore.

`__thread ucontext_t vcore_context`

Context associated with each vcore. Serves as the entry point to this vcore whenever the vcore is first brought up, a usercontext yields on it, or a signal / async I/O notification is to be handled.

`__thread ucontext_t *vcore_saved_ucontext`

Current user context running on each vcore, used when interrupting a user context because of async I/O or signal handling. Vcore 0's `vcore_saved_ucontext` is initialized to the continuation of the main thread's context the first time it's `vcore_entry()` function is invoked.

`__thread void *vcore_saved_tls_desc`

Current `tls_desc` of the user context running on each vcore, used when interrupting a user context because of async I/O or signal handling. Hard Thread 0's `vcore_saved_tls_desc` is initialized to the `tls_desc` of the main thread's context the first time it's `vcore_entry()` function is invoked.

2.1.5 API calls

`int vcore_lib_init();`

Initialization routine for the vcore subsystem.

`void vcore_reenter (void (*entry_func)(void))`

Function to reenter a vcore at the top of its stack.

`int vcore_request (int k)`

Requests `k` additional vcores. Returns -1 if the request is impossible. Otherwise, blocks the calling vcore until the request is granted and returns 0.

`void vcore_yield();`

Relinquishes the calling vcore.

`int vcore_id (void)`

Returns the id of the calling vcore.

`size_t num_vcores (void)`

Returns the current number of vcores allocated.

`size_t max_vcores (void)`

Returns the maximum number of allocatable vcores.

`bool in_vcore_context () {`

Returns whether you are currently running in vcore context or not.

`void clear_notif_pending(uint32_t vcoreid);`

Clears the flag for pending notifications

`void enable_notifs(uint32_t vcoreid);`

Enable Notifications

`void disable_notifs(uint32_t vcoreid);`

Disable Notifications

`#define vcore_begin_access_tls_vars (vcoreid)`

Begin accessing TLS variables associated with a specific vcore (possibly a different from the current one). Matched one-to-one with a following call to `vcore_end_access_tls_vars()` within the same function.


```
#define vcore_end_access_tls_vars ()
    End access to a vcore's TLS variables. Matched one-to-one with a previous call to vcore_begin_access_tls_vars
    () within the same function.

#define vcore_set_tls_var (name, val)
    Set a single variable in the TLS of the current vcore. Mostly useful when running in uthread context and want
    to set something vcore specific.

#define vcore_get_tls_var (name)
    Get a single variable from the TLS of the current vcore. Mostly useful when running in uthread context and
    want to get something vcore specific.
```

2.2 Uthreads

To access the uthread API, include the following header file:

```
#include <parlib/uthread.h>
```

2.2.1 Constants

UTH_EXT_BLK_MUTEX

One, of possibly many in the future, reasons that a uthread has blocked externally. This is required for proper implementation of the **uthread_has_blocked()** API call.

2.2.2 Types

struct uthread

uthread_t

An opaque type used to reference and manage a user-level thread

struct schedule_ops

schedule_ops_t

A struct containing the list of callbacks a user-level scheduler built on top of this uthread library needs to implement.

```
struct schedule_ops {
    void (*sched_entry) (void);
    void (*thread_runnable) (struct uthread *);
    void (*thread_paused) (struct uthread *);
    void (*thread_blockon_sysc) (struct uthread *, void *);
    void (*thread_has_blocked) (struct uthread *, int);
    void (*preempt_pending) (void);
    void (*spawn_thread) (uintptr_t pc_start, void *data);
};
```

void **schedule_ops_t.sched_entry** ()

void **schedule_ops_t.thread_runnable** (struct **uthread** *)

void **schedule_ops_t.thread_paused** (struct **uthread** *)

void **schedule_ops_t.thread_blockon_sysc** (struct **uthread** *, void *)

void **schedule_ops_t.thread_has_blocked** (struct **uthread** *, int)

void **schedule_ops_t.preempt_pending** ()

```
void schedule_ops_t.spawn_thread (uintptr_t pc_start, void *data)
```

2.2.3 External Symbols

```
extern struct schedule_ops *sched_ops
```

A reference to an externally defined variable which contains pointers to implementations of all the `schedule_ops` callbacks.

2.2.4 Global Variables

```
__thread uthread_t *current_uthread
```

2.2.5 API calls

```
int uthread_lib_init ()
```

```
void uthread_cleanup (struct uthread *uthread)
```

```
void uthread_runnable (struct uthread *uthread)
```

```
void uthread_yield (bool save_state, void (*yield_func)(struct uthread*, void*), void *yield_arg)
```

```
void save_current_uthread (struct uthread *uthread)
```

```
void highjack_current_uthread (struct uthread *uthread)
```

```
void run_current_uthread (void)
```

This function does not return.

```
void run_uthread (struct uthread *uthread)
```

This function does not return.

```
void swap_uthreads (struct uthread *__old, struct uthread *__new)
```

```
init_uthread_tf (uthread_t *uth, void (*entry)(void), void *stack_bottom, uint32_t size)
```

```
#define uthread_begin_access_tls_vars (uthread)
```

```
#define uthread_end_access_tls_vars ()
```

```
#define uthread_set_tls_var (uthread, name, val)
```

```
#define uthread_get_tls_var (uthread, name)
```

2.3 MCS Locks

MCS locks are a spinlock style lock designed for more efficient execution on multicore processors. They are designed to mitigate cache clobbering and TLB shootdowns by having each call site spin on a core-local lock variable, rather than the single lock variable used by traditional spinlocks. These locks should really only be used while running in vcore context.

To access the mcs lock API, include the following header file:

```
#include <parlib/mcs.h>
```

2.3.1 Constants

MCS_LOCK_INIT

Static initializer for an `mcs_lock_t`

MCS_QNODE_INIT

Static initializer for an `mcs_lock_qnode_t`

2.3.2 Types

struct mcs_lock_qnode

mcs_lock_qnode_t

An MCS lock qnode. MCS locks are maintained as a queue of MCS qnode pointers, and locks are granted in order of request, using the qnode pointer passed at each `mcs_lock_lock()` call.

struct mcs_lock

mcs_lock_t

An MCS lock itself. This data type keeps track of whether the lock is currently held or not, as well as the list of qnode pointers described above.

struct mcs_dissem_flags

mcs_dissem_flags_t

Dissemination flags used by the MCS barriers described below. This data type should never normally be accessed by an external library. They are used internally by the MCS barrier implementation, but exist as part of the API because the `mcs_barrier` struct contains them.

struct mcs_barrier

mcs_barrier_t

An MCS barrier. This data type is used to reference an MCS barrier across the various MCS barrier API calls.

2.3.3 API calls

void **mcs_lock_init** (struct `mcs_lock` **lock*)

Initializes an MCS lock.

void **mcs_lock_lock** (struct `mcs_lock` **lock*, struct `mcs_lock_qnode` **qnode*)

Locks an MCS lock, associating a call-site specific qnode with the lock in the process.

void **mcs_lock_unlock** (struct `mcs_lock` **lock*, struct `mcs_lock_qnode` **qnode*)

Unlocks an MCS lock, releasing the call-site specific qnode associated with the current lock holder.

void **mcs_lock_notifsafe** (struct `mcs_lock` **lock*, struct `mcs_lock_qnode` **qnode*)

A signal-safe implementation of `mcs_lock_lock()`. While the lock is held, the lockholder will not be interrupted to run any signal handlers.

void **mcs_unlock_notifsafe** (struct `mcs_lock` **lock*, struct `mcs_lock_qnode` **qnode*)

The `mcs_lock_unlock()` counterpart to `mcs_lock_notifsafe()`. After releasing the lock, signals may be processed again.

void **mcs_barrier_init** (`mcs_barrier_t`* *b*, `size_t` *num_vcores*)

Initializes an MCS barrier with the number of vcores associated with the barrier.

void **mcs_barrier_wait** (`mcs_barrier_t`* *b*, `size_t` *vcoid*)

Waits on an MCS barrier for the specified vcoreid.

2.4 Spinlocks

To access the spinlock API, include the following header file:

```
#include <parlib/spinlock.h>
```

2.4.1 Constants

2.4.2 Types

```
struct spinlock  
spinlock_t
```

2.4.3 API calls

```
void spinlock_init (spinlock_t *lock)  
int spinlock_trylock (spinlock_t *lock)  
void spinlock_lock (spinlock_t *lock)  
void spinlock_unlock (spinlock_t *lock)
```

2.5 Dynamic Thread Local Storage

Dynamic thread local storage is parlib's way of providing `pthread_key_create()` and `pthread_set/get_specific()` style semantics for uthreads and vcores. Using dtls, different libraries can dynamically define their own set of dtls keys for thread local storage. Uthreads and vcores can then access the private regions associated with these keys through the API described below.

To access the dynamic thread local storage API, include the following header file:

```
#include <parlib/dtls.h>
```

2.5.1 Types

```
struct dtls_key  
dtls_key_t
```

A dynamic thread local storage key. Uthreads and vcores use these keys to gain access to their own thread local storage region associated with the key. Multiple keys can be created, with each key referring to a different set of thread local storage regions.

```
void (*dtls_dtor_t) (void*);  
dtls_dtor_t
```

A function pointer defining a destructor function mapped to a specific `dtls_key`. Whenever a uthread has accessed a `dtls_key` that has a `dtls_dtor_t` mapped to it, the destructor function will be called before the uthread exits.

2.5.2 API calls

dtls_key_t dtls_key_create (dtls_dtor_t dtor)

Initialize a dtls key for dynamically setting/getting thread local storage on a uthread or vcore. Takes a dtls_dtor_t as a parameter and associates it with a new dtls_key_t, which gets returned.

void dtls_key_delete (dtls_key_t key)

Delete the provided dtls key.

void set_dtls (dtls_key_t key, void *dtls)

Associate a dtls storage region for the provided dtls key on the current uthread or vcore.

void *get_dtls (dtls_key_t key)

Get the dtls storage region previously associated with the provided dtls key on the current uthread or vcore. If no storage region has been associated yet, return NULL.

void destroy_dtls ()

Destroy all dtls storage associated with all keys for the current uthread or vcore.

2.6 Thread Local Storage

To access the thread local storage API, include the following header file:

```
#include <parlib/tls.h>
```

2.6.1 Global Variables

void *main_tls_desc

__thread void *current_tls_desc

2.6.2 API calls

void *allocate_tls (void)

void *reinit_tls (void *tcb)

void free_tls (void *tcb)

void set_tls_desc (void *tls_desc, uint32_t vcoreid)

void *get_tls_desc (uint32_t vcoreid)

2.7 Memory Pools

Memory pools are designed for efficient dynamic memory management when you have a fixed number of fixed size objects you need to manage. A large chunk of memory must first be allocated by traditional means before passing it to the pool API. Once received however, the memory is managed according to the restrictions described above, without worrying about fragmentation or other traditional dynamic memory management concerns.

To access the pool API, include the following header file:

```
#include <parlib/pool.h>
```

2.7.1 Types

struct **pool**
pool_t
Opaque type used to reference and manage a pool

2.7.2 API calls

void **pool_init** (**pool_t** **pool*, void* *buffer*, void ***object_queue*, size_t *num_objects*, size_t *object_size*)
Initialize a pool. All memory MUST be allocated externally. The pool implementation simply manages the objects contained in the buffer passed to this function. This allows us to use the same pool implementation for both statically and dynamically allocated memory.

size_t **pool_size** (**pool_t** **pool*)
Check how many objects the pool is able to hold

size_t **pool_available** (**pool_t** **pool*)
See how many objects are currently available for allocation from the pool.

void* **pool_alloc** (**pool_t** **pool*)
Get an object from the pool

int **pool_free** (**pool_t** **pool*, void **object*)
Put an object into the pool

2.8 Slab Memory Allocator

To access the slab allocator API, include the following header file:

```
#include <parlib/slab.h>
```

2.8.1 Types

struct **slab_cache**
slab_cache_t

void (***slab_ctor_t**) (void *, size_t)
slab_ctor_t

void (***slab_dtor_t**) (void *, size_t)
slab_dtor_t

2.8.2 API calls

struct **slab_cache** ***slab_cache_create** (const char **name*, size_t *obj_size*, int *align*, int *flags*,
slab_ctor_t *ctor*, slab_dtor_t *dtor*)

void **slab_cache_destroy** (struct **slab_cache** **cp*)

void ***slab_cache_alloc** (struct **slab_cache** **cp*, int *flags*)

void **slab_cache_free** (struct **slab_cache** **cp*, void **buf*)

2.9 Atomic Memory Operations

To access the atomic memory operations API, include the following header file:

```
#include <parlib/atomic.h>
```

2.9.1 Types

atomic_t

2.9.2 API calls

```
void atomic_init (atomic_t *number, long val)
void *atomic_swap_ptr (void **addr, void *val)
long atomic_swap (atomic_t *addr, long val)
uint32_t atomic_swap_u32 (uint32_t *addr, uint32_t val)
#define mb ()
#define cmb ()
#define rmb ()
#define wmb ()
#define wrmb ()
#define rwmb ()
#define mb_f ()
#define rmb_f ()
#define wmb_f ()
#define wrmb_f ()
#define rwmb_f ()
```


ABOUT

3.1 People

Current Contributors:

- Kevin Klues <klueska@cs.berkeley.edu>
- Barret Rhoden <brho@cs.berkeley.edu>

3.2 Publications

There are currently no publications on Parlib.

INDEX

Symbols

__vcore_tls_descs (C variable), 4
__vcores (C variable), 4

A

allocate_tls (C function), 9
atomic_init (C function), 11
atomic_swap (C function), 11
atomic_swap_ptr (C function), 11
atomic_swap_u32 (C function), 11
atomic_t (C type), 11

C

cmb (C function), 11
current_tls_desc (C variable), 9
current_uthread (C variable), 6

D

destroy_dtls (C function), 9
dtls_dtor_t (C type), 8
dtls_key (C type), 8
dtls_key_create (C function), 9
dtls_key_delete (C function), 9
dtls_key_t (C type), 8

F

free_tls (C function), 9

G

get_dtls (C function), 9
get_tls_desc (C function), 9

H

highjack_current_uthread (C function), 6

I

init_uthread_tf (C function), 6

L

LOG2_MAX_VCORES (C macro), 3

M

main_tls_desc (C variable), 9
max_vcores (C function), 4
MAX_VCORES (C macro), 3
mb (C function), 11
mb_f (C function), 11
mcs_barrier (C type), 7
mcs_barrier_init (C function), 7
mcs_barrier_t (C type), 7
mcs_barrier_wait (C function), 7
mcs_dissem_flags (C type), 7
mcs_dissem_flags_t (C type), 7
mcs_lock (C type), 7
mcs_lock_init (C function), 7
MCS_LOCK_INIT (C macro), 7
mcs_lock_lock (C function), 7
mcs_lock_notifsafe (C function), 7
mcs_lock_qnode (C type), 7
mcs_lock_qnode_t (C type), 7
mcs_lock_t (C type), 7
mcs_lock_unlock (C function), 7
MCS_QNODE_INIT (C macro), 7
mcs_unlock_notifsafe (C function), 7

N

num_vcores (C function), 4

P

pool (C type), 10
pool_alloc (C function), 10
pool_available (C function), 10
pool_free (C function), 10
pool_init (C function), 10
pool_size (C function), 10
pool_t (C type), 10

R

reinit_tls (C function), 9
rmb (C function), 11
rmb_f (C function), 11
run_current_uthread (C function), 6

run_uthread (C function), 6
rwmb (C function), 11
rwmb_f (C function), 11

S

save_current_uthread (C function), 6
sched_ops (C variable), 6
schedule_ops (C type), 5
schedule_ops_t (C type), 5
schedule_ops_t.preempt_pending (C function), 5
schedule_ops_t.sched_entry (C function), 5
schedule_ops_t.spawn_thread (C function), 5
schedule_ops_t.thread_blockon_sysc (C function), 5
schedule_ops_t.thread_has_blocked (C function), 5
schedule_ops_t.thread_paused (C function), 5
schedule_ops_t.thread_runnable (C function), 5
set_dtls (C function), 9
set_tls_desc (C function), 9
slab_cache (C type), 10
slab_cache_alloc (C function), 10
slab_cache_create (C function), 10
slab_cache_destroy (C function), 10
slab_cache_free (C function), 10
slab_cache_t (C type), 10
slab_ctor_t (C type), 10
slab_dtor_t (C type), 10
spinlock (C type), 8
spinlock_init (C function), 8
spinlock_lock (C function), 8
spinlock_t (C type), 8
spinlock_trylock (C function), 8
spinlock_unlock (C function), 8
swap_uthreads (C function), 6

U

UTH_EXT_BLK_MUTEX (C macro), 5
uthread (C type), 5
uthread_begin_access_tls_vars (C function), 6
uthread_cleanup (C function), 6
uthread_end_access_tls_vars (C function), 6
uthread_get_tls_var (C function), 6
uthread_lib_init (C function), 6
uthread_runnable (C function), 6
uthread_set_tls_var (C function), 6
uthread_t (C type), 5
uthread_yield (C function), 6

V

vcore (C type), 3
vcore_begin_access_tls_vars (C function), 4
vcore_context (C variable), 4
vcore_end_access_tls_vars (C function), 5
vcore_entry (C function), 3
vcore_get_tls_var (C function), 5

vcore_id (C function), 4
vcore_reenter (C function), 4
vcore_request (C function), 4
vcore_saved_tls_desc (C variable), 4
vcore_saved_ucontext (C variable), 4
vcore_set_tls_var (C function), 5
vcore_t (C type), 3

W

wmb (C function), 11
wmb_f (C function), 11
wrmb (C function), 11
wrmb_f (C function), 11